
Bayesian-Adaptive Deep Reinforcement Learning via Ensemble Learning

Gilwoo Lee
gilwoo@cs.uw.edu

Jeongseok Lee
jslee02@cs.uw.edu

1 Introduction

While reinforcement learning is capable of controlling complex autonomous systems, RL algorithms typically require huge amounts of data and can overfit to a particular task or to be prone to disturbance. One of main challenges that needs to be addressed is train a policy robust to various model uncertainties and disturbances. In this project, we aim to address this challenge via an ensemble policy for Bayes-Adaptive Reinforcement Learning [1].

We assume that there exists a latent physics variable ϕ which determines the transition function of the underlying MDP, i.e. the transition function $P(s', \phi' | s, \phi, a)$ is now a function of state, action, and ϕ . We would like to learn a policy which maximizes the long term reward given ϕ . Formally, this is called Bayes-Adaptive MDP [1, 2], defined by a tuple $\langle \mathcal{S}', \mathcal{A}, P', P_0, R' \rangle$ where

- $\mathcal{S}' = \mathcal{S} \times \Phi$ is the set of (states, physics variable),
- \mathcal{A} is the set of actions,
- $P(\cdot | s, \phi, a)$ is the transition function between hyper-states, conditioned on action a being taken in hyper-state (s, ϕ) ,
- $P_0 \in \mathcal{P}(\mathcal{S} \times \Phi)$ combines the initial distribution over hyper-states,
- $R'(s, \phi, a)$ represents the reward obtained when action a is taken in hyper-state (s, ϕ) .

We would like to find the optimal policy for the following Bellman equation:

$$V^*(s, \phi) = \max_a \mathbb{E} \left[R(s, a, \phi) + \gamma \sum_{s', \phi'} P(s', \phi' | s, \phi, a) V^*(s', \phi') \right] \quad (1)$$

This formulation is often referred as Bayes-Adaptive Reinforcement Learning (BARL) [1].

We make a simplification to BARL formulation. We assume that the dynamics of s' and ϕ' are independent given $P(s, \phi, a)$, i.e.

$$P(s', \phi' | s, \phi, a) = P(s' | s, \phi, a) \cdot P(\phi' | s, \phi, a).$$

This assumption allows us to simplify BARL with a gated ensemble policy learning method. At the high-level, we have a gating network that determines the *belief* of the physics parameters at time t ,

$$b(\phi_t) = P(\phi_t | s_{t-1}, \phi_{t-1}, a_{t-1})$$

which is then used to compute the best policy from an ensemble of ϕ -dependent optimal policies, i.e., $\pi^*(\cdot; \phi)$ and $V^*(\cdot; \phi)$ are computed with typical RL algorithms for MDPs. Then the remaining task is to compute the one-step best action a :

$$a^* = \arg \max_a \mathbb{E}_{\phi \sim b(\phi)} \left[R(s, a, \phi) + \gamma \sum_{s', \phi'} P(s', \phi' | s, \phi, a) V^*(s', \phi') \right] \quad (2)$$

We model b_ϕ as a network capable of modeling evolving state change, e.g. Recurrent Neural Networks, or as a Bayes filter. At the low level, we plan to discretize Φ and have one actor-critic network per a discretized value of ϕ : each critic estimates $V^*(\cdot; \phi)$ and each actor has an optimal policy for a particular discretized value of $\pi^*(\cdot; \phi)$. Given b_ϕ and the set of actor-critic networks, it is straightforward to compute (2).

2 Background

Our work is closely related to QMDP [3, 4] which is an approximation for POMDP. QMDP approximates POMDP by assuming fully-observable MDP after 1-step, and approximating the Q-value at the current belief state $b(s)$ as $Q_a(b) = \sum_s b(s)Q_{MDP}(s, a)$. In our problem setup, we have a belief over the physics parameters ϕ of the MDP, $b(\phi)$, and we compute the policy $Q_a(s; b) = \sum_{\phi} b(\phi)Q_{MDP}(s, a; \phi)$.

The BAMDP formulation is also similar to POMDP formulation used in POMDP-lite [5] which assumes that the hidden state variables are constant or only change deterministically. In our case, the hidden state variables correspond to the physics parameters ϕ . The authors of POMDP-lite have shown that such formulation is “equivalent to a set of fully observable Markov decision processes indexed by a hidden parameter” [5], which, in our case, could be viewed as a discretization of ϕ .

3 Platform

We plan to simulate with a set of canonical examples such as inverted pendulum, 2D car driving, hopper, and do a real-robot experiments with RACECAR, a RC-car robot actively developed in Personal Robotics Lab.

4 Milestones

References

- [1] M. Ghavamzadeh, S. Mannor, J. Pineau, A. Tamar, *et al.*, “Bayesian reinforcement learning: A survey,” *Foundations and Trends® in Machine Learning*, vol. 8, no. 5-6, pp. 359–483, 2015.
- [2] A. Guez, D. Silver, and P. Dayan, “Efficient bayes-adaptive reinforcement learning using sample-based search,” in *Advances in Neural Information Processing Systems*, pp. 1025–1033, 2012.
- [3] M. L. Littman, A. R. Cassandra, and L. P. Kaelbling, “Learning policies for partially observable environments: Scaling up,” in *Machine Learning Proceedings 1995*, pp. 362–370, Elsevier, 1995.
- [4] P. Karkus, D. Hsu, and W. S. Lee, “Qmdp-net: Deep learning for planning under partial observability,” in *Advances in Neural Information Processing Systems*, pp. 4697–4707, 2017.
- [5] M. Chen, E. Frazzoli, D. Hsu, and W. S. Lee, “Pomdp-lite for robust robot planning under uncertainty,” in *Robotics and Automation (ICRA), 2016 IEEE International Conference on*, pp. 5427–5433, IEEE, 2016.